Social Science Inquiry II

Week 9: Beyond linear regression, part I

Molly Offer-Westort

Department of Political Science,
University of Chicago

Winter 2022

Loading packages for this class

```
> set.seed(60637)
> library(ggplot2)
```

- Housekeeping

Machine learning

What is it?

Machine learning

What is it?

- A body of *algorithmic* methods . . .

Machine learning

What is it?

- A body of *algorithmic* methods . . . (an algorithm is just a recipe)

Machine learning

What is it?

- A body of *algorithmic* methods ... (an algorithm is just a recipe)
- Somehow part of *artificial intelligence*...

Machine learning

What is it?

- A body of *algorithmic* methods . . . (an algorithm is just a recipe)
- Somehow part of *artificial intelligence*. . . (basically, how computers perform tasks)

Machine learning

What is it?

- A body of *algorithmic* methods . . . (an algorithm is just a recipe)
- Somehow part of *artificial intelligence*. . . (basically, how computers perform tasks)
- In general, a flexible, *data-driven* approach to make predictions, classify data, or take decisions.

Machine learning

How do the objectives of machine learning differ from those of conventional quantitative methods in the social sciences?

Machine learning

How do the objectives of machine learning differ from those of conventional quantitative methods in the social sciences?

- In conventional quantitative methods:
  - identify and estimate a target estimand, which is often a parameter in a statistical model,

Machine learning

How do the objectives of machine learning differ from those of conventional quantitative methods in the social sciences?

- In conventional quantitative methods:
  - identify and estimate a target estimand, which is often a parameter in a statistical model, defined over some specified population of interest.

Machine learning

How do the objectives of machine learning differ from those of conventional quantitative methods in the social sciences?

- In conventional quantitative methods:
  - identify and estimate a target estimand, which is often a parameter in a statistical model, defined over some specified population of interest.
  - descriptive or causal, e.g., population employment rate, or effect of a policy

Machine learning

How do the objectives of machine learning differ from those of conventional quantitative methods in the social sciences?

- In conventional quantitative methods:
  - identify and estimate a target estimand, which is often a parameter in a statistical model, defined over some specified population of interest.
  - descriptive or causal, e.g., population employment rate, or effect of a policy
- In ML:
  - development of algorithms to make classifications or predictions.

Machine learning

How do the objectives of machine learning differ from those of conventional quantitative methods in the social sciences?

- In conventional quantitative methods:
  - identify and estimate a target estimand, which is often a parameter in a statistical model, defined over some specified population of interest.
  - descriptive or causal, e.g., population employment rate, or effect of a policy
- In ML:
  - development of algorithms to make classifications or predictions.
  - e.g., is this a picture of banana or a cat?

Machine learning

How do the objectives of machine learning differ from those of conventional quantitative methods in the social sciences?

- In conventional quantitative methods:
  - identify and estimate a target estimand, which is often a parameter in a statistical model, defined over some specified population of interest.
  - descriptive or causal, e.g., population employment rate, or effect of a policy
- In ML:
  - development of algorithms to make classifications or predictions.
  - e.g., is this a picture of banana or a cat? Will this person be more likely to click on an ad for sneakers or cookware?

Machine learning

How do the objectives of machine learning differ from those of conventional quantitative methods in the social sciences?

- In conventional quantitative methods:
  - identify and estimate a target estimand, which is often a parameter in a statistical model, defined over some specified population of interest.
  - descriptive or causal, e.g., population employment rate, or effect of a policy
- In ML:
  - development of algorithms to make classifications or predictions.
  - e.g., is this a picture of banana or a cat? Will this person be more likely to click on an ad for sneakers or cookware?
- Is there overlap between the two?

Model fit vs. prediction

- In linear regression, propose a model:

$$\hat{Y}_i = \hat{\beta}_0 + \hat{\beta}_1 X_{1i} + \hat{\beta}_2 X_{2i} + \cdots + \hat{\beta}_K X_{Ki}$$

- Select $\hat{\beta}_0 \ldots \hat{\beta}_K$ to minimize

$$\sum_{i=1}^{N} \hat{\varepsilon}_i{}^2 = \sum_{i=1}^{N} \left(\hat{Y}_i - Y_i\right)^2$$

Model fit vs. prediction

- For prediction tasks, we could use the same model,

$$\hat{Y}_i = \hat{\beta}_0 + \hat{\beta}_1 X_{1i} + \hat{\beta}_2 X_{2i} + \cdots + \hat{\beta}_K X_{Ki}$$

- But select $\hat{\beta}_0 \ldots \hat{\beta}_K$ to minimize squared prediction error for the next observation:

$$\hat{\varepsilon}^2_{N+1} = \left(\hat{Y}_{N+1} - Y_{N+1}\right)^2$$

Model fit vs. prediction

- For prediction tasks, we could use the same model,

$$\hat{Y}_i = \hat{\beta}_0 + \hat{\beta}_1 X_{1i} + \hat{\beta}_2 X_{2i} + \cdots + \hat{\beta}_K X_{Ki}$$

- But select $\hat{\beta}_0 \ldots \hat{\beta}_K$ to minimize squared prediction error for the next observation:

$$\hat{\varepsilon}^2_{N+1} = \left(\hat{Y}_{N+1} - Y_{N+1}\right)^2$$

- Are these the same thing?

Model fit vs. prediction

- For prediction tasks, we could use the same model,

$$\hat{Y}_i = \hat{\beta}_0 + \hat{\beta}_1 X_{1i} + \hat{\beta}_2 X_{2i} + \cdots + \hat{\beta}_K X_{Ki}$$

- But select $\hat{\beta}_0 \ldots \hat{\beta}_K$ to minimize squared prediction error for the next observation:

$$\hat{\varepsilon}^2_{N+1} = \left(\hat{Y}_{N+1} - Y_{N+1}\right)^2$$

- Are these the same thing? (no)

Model fit vs. prediction

- For prediction tasks, we could use the same model,

$$\hat{Y}_i = \hat{\beta}_0 + \hat{\beta}_1 X_{1i} + \hat{\beta}_2 X_{2i} + \cdots + \hat{\beta}_K X_{Ki}$$

- But select $\hat{\beta}_0 \ldots \hat{\beta}_K$ to minimize squared prediction error for the next observation:

$$\hat{\varepsilon}^2_{N+1} = \left(\hat{Y}_{N+1} - Y_{N+1}\right)^2$$

- Are these the same thing? (no)
- If prediction is our goal, can we do better than least squares regression?

Model fit vs. prediction

- For prediction tasks, we could use the same model,

$$\hat{Y}_i = \hat{\beta}_0 + \hat{\beta}_1 X_{1i} + \hat{\beta}_2 X_{2i} + \cdots + \hat{\beta}_K X_{Ki}$$

- But select $\hat{\beta}_0 \ldots \hat{\beta}_K$ to minimize squared prediction error for the next observation:

$$\hat{\varepsilon}^2_{N+1} = \left(\hat{Y}_{N+1} - Y_{N+1}\right)^2$$

- Are these the same thing? (no)
- If prediction is our goal, can we do better than least squares regression? (yes)

Some ML tools

- A major concern of ML: *overfit*
  - If your model fits the data *too* perfectly, it's not useful for prediction

Suppose we would like to fit a model to the following data:

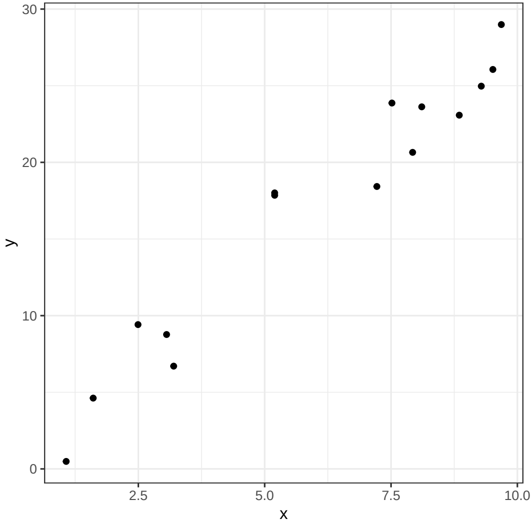

We could use a single line:

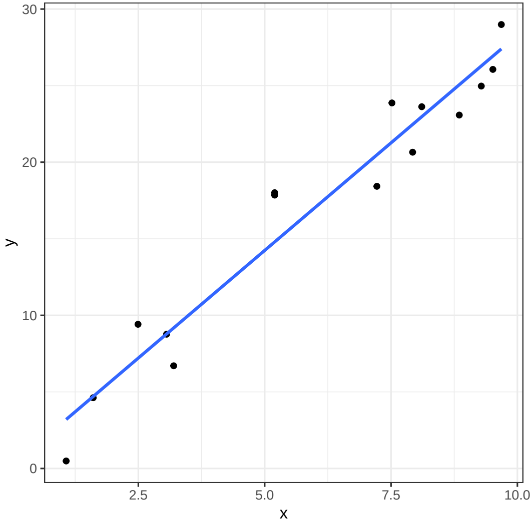

Or we could fit a line between every point:

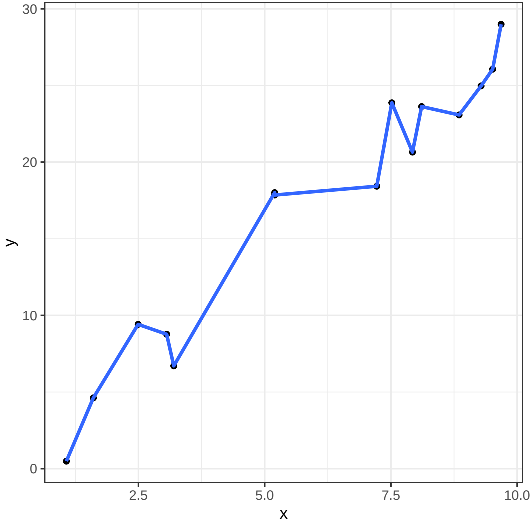

Cross-validation

- If we were to draw another observation from the joint distribution of (Y, X), which one do you think would do a better job of prediction?

Cross-validation

- If we were to draw another observation from the joint distribution of (Y, X), which one do you think would do a better job of prediction?
- ML methods propose a way to check this, by separating data into training and test sets.

Cross-validation

- If we were to draw another observation from the joint distribution of (Y, X), which one do you think would do a better job of prediction?
- ML methods propose a way to check this, by separating data into training and test sets.
- You can fit different models on the training set, and then see which one does the best job of predicting response in the test set.

Cross-validation

- If we were to draw another observation from the joint distribution of (Y, X), which one do you think would do a better job of prediction?
- ML methods propose a way to check this, by separating data into training and test sets.
- You can fit different models on the training set, and then see which one does the best job of predicting response in the test set. (This is not a new idea.)
- There are some different ways to do this:
  - Leave-k-out
  - Leave-one-out
  - k-fold cross validation

Regularization

- Overfit can become a real problem when we have a lot of predictors (K) relative to our number of observations (N)

Regularization

- Overfit can become a real problem when we have a lot of predictors (K) relative to our number of observations (N)
- This is a common problem when we think about an industry setting, where for every customer a business might have a large number of measurements.

Regularization

- Overfit can become a real problem when we have a lot of predictors (K) relative to our number of observations (N)
- This is a common problem when we think about an industry setting, where for every customer a business might have a large number of measurements. Which ones should they use to predict an outcome?

Regularization

- Consider a model:

$$Y = \beta_0 + \beta_1 X_1 + \beta_2 X_2 + \cdots + \beta_K X_K + \varepsilon$$

Regularization

- Consider a model:

$$Y = \beta_0 + \beta_1 X_1 + \beta_2 X_2 + \cdots + \beta_K X_K + \varepsilon$$

- With $K \geq N$, even if every β_k is non-zero, we won't be able to make good predictions with all of our $\hat{\beta}_k$–and when we care about prediction, that's not our goal, anyhow.

Regularization

- Consider a model:

$$Y = \beta_0 + \beta_1 X_1 + \beta_2 X_2 + \cdots + \beta_K X_K + \varepsilon$$

- With $K \geq N$, even if every β_k is non-zero, we won't be able to make good predictions with all of our $\hat{\beta}_k$–and when we care about prediction, that's not our goal, anyhow.
- With regularization, we shrink some of the $\hat{\beta}_k$ nearly all the way or all the way to zero.

Regularization

- Consider a model:

$$Y = \beta_0 + \beta_1 X_1 + \beta_2 X_2 + \cdots + \beta_K X_K + \varepsilon$$

- With $K \geq N$, even if every β_k is non-zero, we won't be able to make good predictions with all of our $\hat{\beta}_k$–and when we care about prediction, that's not our goal, anyhow.
- With regularization, we shrink some of the $\hat{\beta}_k$ nearly all the way or all the way to zero.
- For *ridge regression* or *lasso*, we select the $\hat{\beta}_k$ using:

$$\underset{\boldsymbol{\beta}}{\operatorname{argmin}} \left\{ \sum_{i=1}^{N} \left(Y_i - \beta_0 + \sum_{k=1}^{K} X_{ki} \beta_k \right)^2 + \lambda \sum_{k=1}^{K} |\beta_k|^q \right\}$$

Hastie et al. (2009)

Regression Trees

- Suppose we have joint data, (Y, X), with just one predictor, X.

Regression Trees

- Suppose we have joint data, (Y, X), with just one predictor, X.
- Our goal is to pick some value of c so that we can split the data into two sub-samples ...
  - $X_i \leq c$
  - $X_i > c$

Regression Trees

- Suppose we have joint data, (Y, X), with just one predictor, X.
- Our goal is to pick some value of c so that we can split the data into two sub-samples ...
  - $X_i \leq c$
  - $X_i > c$
- ... and for each sub-sample, predict $\hat{Y}$ as the mean of the Y_i within each sample.

Regression Trees

- Suppose we have joint data, (Y, X), with just one predictor, X.
- Our goal is to pick some value of c so that we can split the data into two sub-samples ...
  - $X_i \leq c$
  - $X_i > c$
- ... and for each sub-sample, predict $\hat{Y}$ as the mean of the Y_i within each sample.
- We want to pick c to minimize:

$$Q = \sum_{i:X_i \leq c} (Y_i - \bar{Y}_{\text{lower}})^2 + \sum_{i:X_i > c} (Y_i - \bar{Y}_{\text{upper}})^2$$

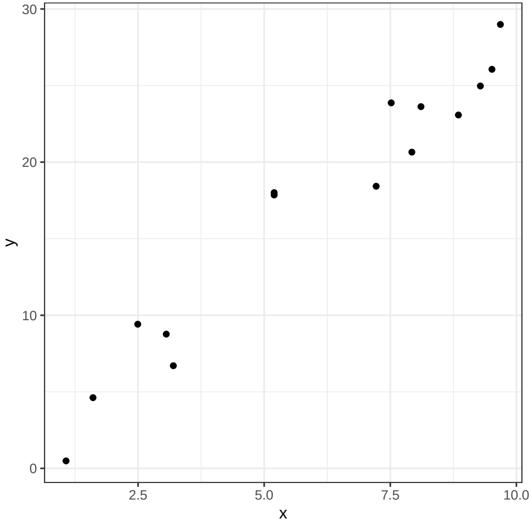
30
20
10
0
y
2.5
5.0
7.5
10.0
x

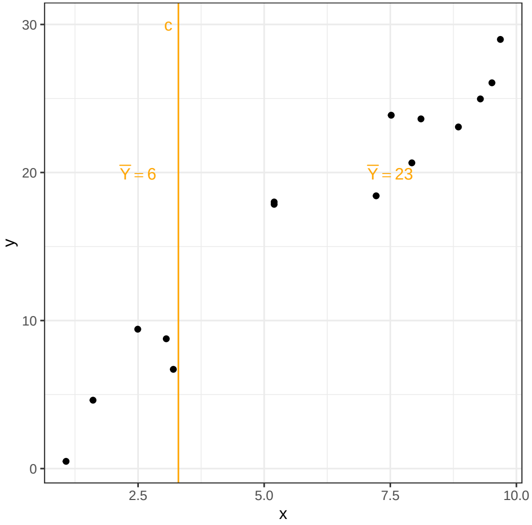
30
20
10
0
2.5
5.0
7.5
10.0
x
y
c
$\overline{Y} = 6$
$\overline{Y} = 23$

Regression Trees

- Now suppose we have joint data, $(Y, X_1, \ldots, X_k)$.

Regression Trees

- Now suppose we have joint data, $(Y, X_1, \ldots, X_k)$.
- We will do the same approach to finding thresholds to minimize prediction error, but we'll want to pick which X_k we use for threshholding, as well.

Regression Trees

- Now suppose we have joint data, $(Y, X_1, \ldots, X_k)$.
- We will do the same approach to finding thresholds to minimize prediction error, but we'll want to pick which X_k we use for threshholding, as well.
- Generally, we'll define the depth of the tree as 2 or three variables; first we'll split on X_k, then we'll split on $X_j \ldots$

Matrix completion: the Nexflix problem

- Netflix has data on viewers, their characteristics, and how they rate movies.

Matrix completion: the Nexflix problem

- Netflix has data on viewers, their characteristics, and how they rate movies.
- The question: how to best recommend to them movies that they have not yet rated?

Matrix completion: the Nexflix problem

- Netflix has data on viewers, their characteristics, and how they rate movies.
- The question: how to best recommend to them movies that they have not yet rated?
- The challenge: come up with the best recommendation algorithm, winner gets $1 million.

Matrix completion: the Nexflix problem

- Netflix has data on viewers, their characteristics, and how they rate movies.
- The question: how to best recommend to them movies that they have not yet rated?
- The challenge: come up with the best recommendation algorithm, winner gets $1 million.
- This can be framed as a matrix completion problem: put users on rows, movies on columns, predict all of the missing rankings.

Causal inference

- Machine learning tools can be super useful for causal inference.

Causal inference

- Machine learning tools can be super useful for causal inference.
  - Fit prediction models separately to treatment and control, so we can do a better job of estimating treatment effects at different covariate values.

Causal inference

- Machine learning tools can be super useful for causal inference.
  - Fit prediction models separately to treatment and control, so we can do a better job of estimating treatment effects at different covariate values.
  - Learn which covariates to include in a (causal) regression model.

Causal inference

- Machine learning tools can be super useful for causal inference.
  - Fit prediction models separately to treatment and control, so we can do a better job of estimating treatment effects at different covariate values.
  - Learn which covariates to include in a (causal) regression model.
  - For observational data, predict propensity to be in treatment vs. control group, based on covariates.

Causal inference: no free lunch

- Machine learning does not solve the fundamental problem of causal inference.

Causal inference: no free lunch

- Machine learning does not solve the fundamental problem of causal inference.
- Causal interpretations are based on assumptions about the data generating process, or knowledge of assignment procedures. These are outside the realm of machine learning methods.

Prediction error

- ML methods may do a good job of producing estimates, but how do we account for inference?

Prediction error

- ML methods may do a good job of producing estimates, but how do we account for inference?
- Cross-validation

Prediction error

- ML methods may do a good job of producing estimates, but how do we account for inference?
- Cross-validation
- Bootstrapping

Prediction error

- ML methods may do a good job of producing estimates, but how do we account for inference?
- Cross-validation
- Bootstrapping
- Applying these solutions to prediction under multiple linear regression

References I

Athey, S. and Imbens, G. W. (2019). Machine learning methods that economists should know about. Annual Review of Economics, 11:685–725.

Hastie, T., Tibshirani, R., Friedman, J. H., and Friedman, J. H. (2009). The elements of statistical learning: data mining, inference, and prediction, volume 2. Springer.